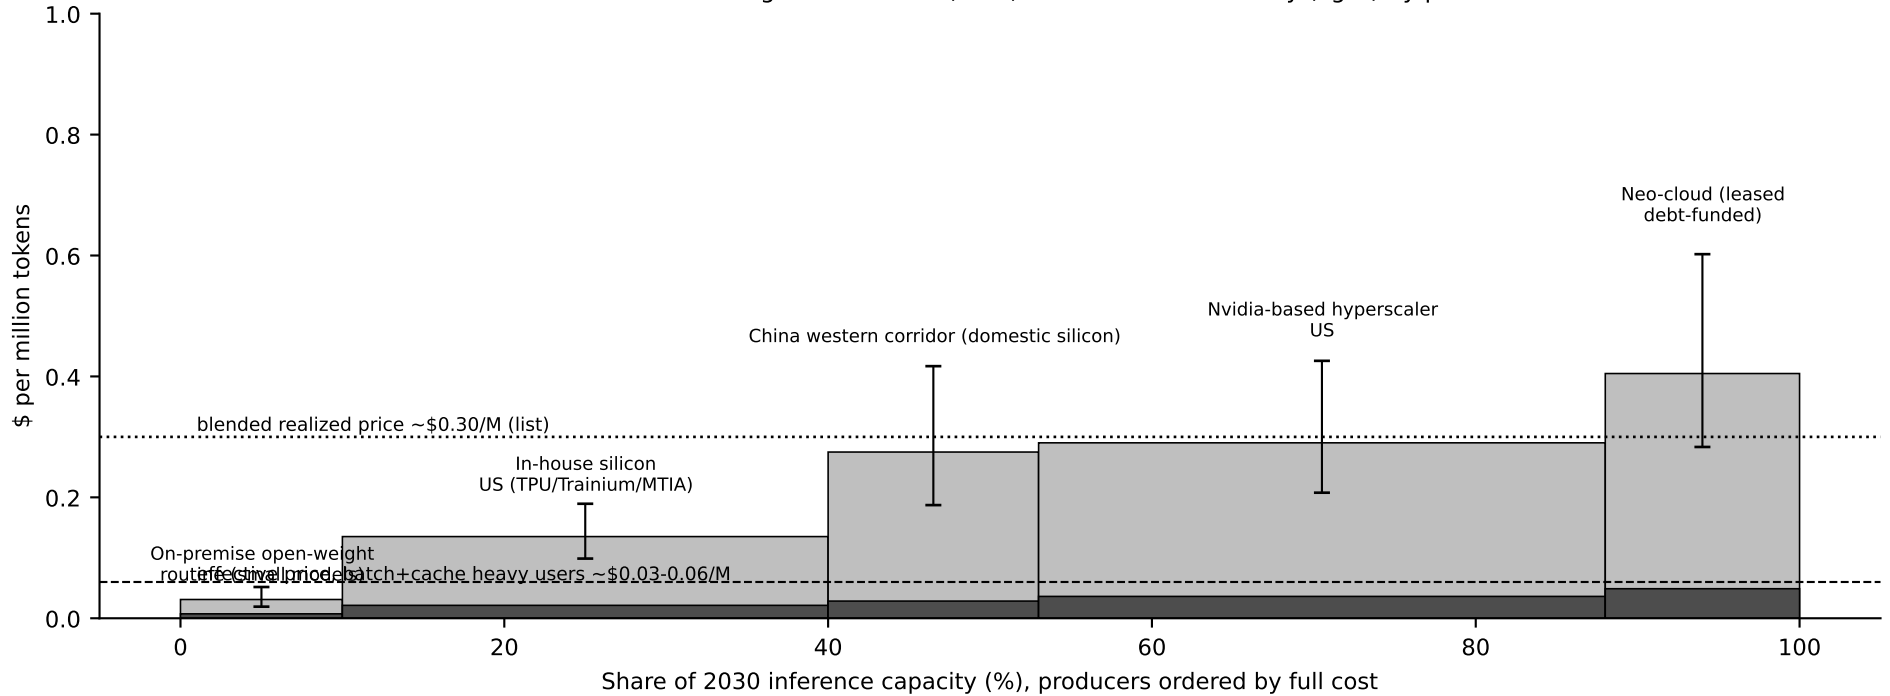
The inference cost curve: marginal cash cost (dark) and full cost incl. carry (light) by producer tier
$ per million tokens
1.0
0.8
0.6
0.4
0.2
0.0
0
20
40
60
80
100
Share of 2030 inference capacity (%), producers ordered by full cost
blended realized price ~$0.30/M (list)
effective price
On-premise open-weight
router + batch+cache heavy users ~$0.03-0.06/M
In-house silicon
US (TPU/Trainium/MTIA)
China western corridor (domestic silicon)
Nvidia-based hyperscaler
US
Neo-cloud (leased
debt-funded)