Will the AI build-out pay for itself?

The argument in twelve slides

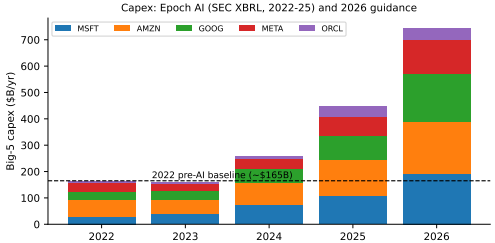
Desmond Fitzpatrick Retired, IBM Systems Founder, Nvara.ai
desmond.fitzpatrick@nvara.ai desmond.fitzpatrick@gmail.com

This is the short version of a 105-page paper whose every number sits in a public register that anyone can challenge. If a slide here and the paper disagree, the paper is right.

Paper, model, data and the register: see the last slide.

1. The question, in one sentence

Five companies — Microsoft, Alphabet, Amazon, Meta and Oracle — are spending about **$780 billion this year** on AI data centers, after $448 billion last year. The question is not whether AI is useful. It is whether the *revenue* will arrive, fast enough, to pay for the buildings and the chips.



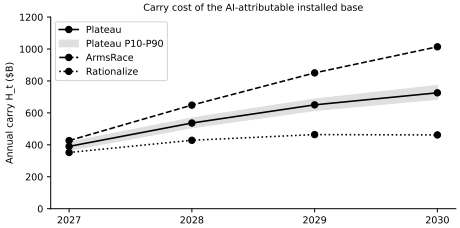
Capex from company filings via Epoch AI; 2026 from the companies' own July guidance. Street consensus for 2027 is higher still, about $1.05 trillion.

2. What the build-out costs every year, not just once

A data center is not paid for when it is bought. Each year it costs its owner three things: the chips wearing out (they last about four years), running the building (power, people, maintenance), and the return the money could have earned elsewhere. We call that yearly cost the **carry**.

$390B → $725B

carry per year, 2027 to 2030, if spending merely stays flat



Why it keeps rising even if spending stops rising: this year's chips are still on the books when next year's arrive. The pile grows faster than it wears out.

Paper Section 4. About 28 cents of revenue per dollar of capital per year; Goldman Sachs's September 2026 framework says 27.

3. What comes in today

All AI revenue in calendar 2026 — the frontier labs (OpenAI, Anthropic and others), the clouds' AI services, the AI lift in advertising, consumer subscriptions — adds up to roughly **$180 billion**. That is real money and it is growing very fast. It is also less than a quarter of what the installed base will cost to carry in 2030.

~$180B

AI revenue, 2026, all channels (range $165–190B)

~$725B

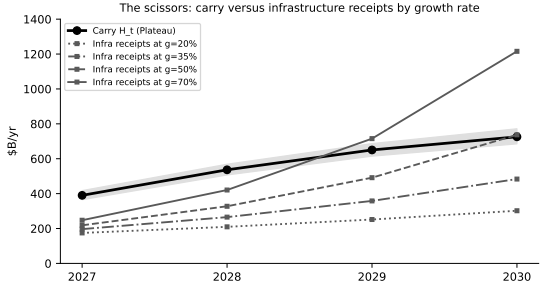
carry of the installed base, 2030, if spending plateaus

So the whole argument is about one number: **how fast revenue has to grow** to catch the carry — and whether anything in the real economy can grow that fast for that long.

Paper Section 7: the 2026 base, decomposed into seven sourced pieces; lab run-rates are discounted to booked revenue.

4. The scissors: how fast revenue must grow

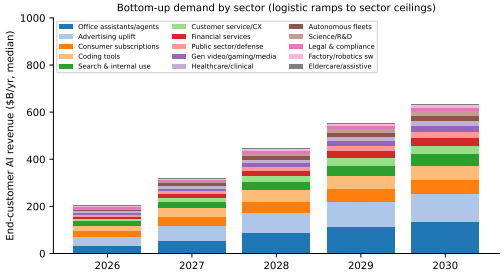
Grow today's $180 billion at a steady rate and see where it crosses the carry. It takes about **49% a year, every year to 2030**, just to break even on a flat spending path — and nothing clears 2027–28, which is already a sunk loss of $100–250 billion a year.



Paper Section 10. Under the street's own spending path (an arms race) break-even needs 62% a year; under a 35% cut, 33%.

5. Where would that revenue come from? We counted, sector by sector

Instead of assuming a growth rate, the paper builds demand up from fifteen sectors — office assistants and agents, advertising, consumer apps, coding tools, healthcare, robotics, and so on — each with a 2026 base, a ceiling set by payrolls, IT budgets or users, and a realistic ramp.



~$625B

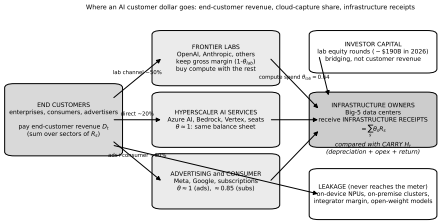
what end customers spend on AI in 2030 (range $525–745B)

That is about **33% a year**, not 49%. Growth starts near 57% in 2027 and decays toward 14% by 2030 as each sector nears its ceiling.

Paper Section 8. No single sector closes the gap; the four largest are office assistants, advertising, consumer and coding.

6. And not all of it reaches the people who built the data centers

A customer's dollar passes through the labs and the application builders before any of it reaches the cloud owner. Routine work moves onto cheap or local chips; outcome-priced integrators keep most of the savings they create. In contact centers, roughly $50 of payroll disappears for every $1 of cloud tokens.



Paper Sections 8 and 13. Of the $625B customers spend in 2030, about $480–490B reaches the infrastructure owners.

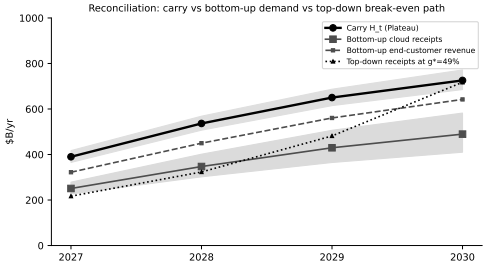
7. The gap

~$255B a year

short in 2030 if spending plateaus

$190–540B

short in 2030 under the street's arms-race path



This is 3.5 standard deviations of everything we do not know, with all the ranges propagated honestly. It is not a rounding error.

Paper Section 10, Table 12. 2027–28 is a loss on every path; the argument is about whether 2029–30 recovers it.

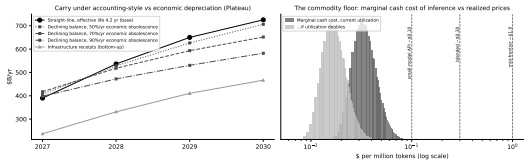
8. Why we think growth slows: five things that are already visible

1. **Slack.** Saving an analyst thirty minutes a day across six tasks does not let you fire anyone. Nominal time savings convert to cash at roughly 28 cents on the dollar.
2. **Verification.** The person who used to write the memo now checks it. Liability keeps a human in the loop.
3. **Leakage.** Savings go to the buyer and the integrator who fixed the process, not the cloud. Routine inference moves to 1.45 billion on-device chips a year.
4. **Sovereignty.** Europe, India and others carve 15–35% of demand out of the US clouds while forcing them to duplicate capital.
5. **Grid and consent.** $130 billion of US projects were blocked or delayed in the first quarter of 2026 alone. Chips bought before they can be powered earn nothing.

Paper Sections 11–13. Each has a precedent: electricity took forty years to show up in factory productivity; computers took twenty.

9. The chips are already worth less than the books say

Accountants depreciate a chip over five or six years. Economically it is worth what it can earn against the newest chip, whose cost per token falls 50–90% a year. On that schedule the silicon bought through 2026 is already worth **$100–250 billion less** than the balance sheets show.



Paper Section 14. If token prices fall to the cash cost of running the chips (a tenth of today's), the capital is never recovered.

10. What the strongest bull cases say — and where we actually disagree

The paper takes a16z's September 2026 *State of Markets* and Goldman Sachs's September 2026 ROIC note point by point. We agree with both on almost every observable: capex consensus, backlogs near $1.7 trillion, rising GPU rents, the ROIC on existing businesses. The yearly cost per dollar of capital is not in dispute — Goldman's 27 cents, our 28.

The disagreements are three, and they are judgments, not arithmetic:

- Whether AI gets paid from **labor budgets** (the $30 trillion payroll) as well as IT budgets — the bulls say yes; we show why only a fraction converts.
- Whether today's scarce, expensive compute is the **permanent state** or the top of a cycle.
- Whether you count **one cohort** of spending (Goldman's $1.73 trillion for 2026–27) or every cohort stacked on the balance sheet. A balance sheet does not get to choose.

Paper Sections 15 and 16.

11. What would prove this wrong — and when

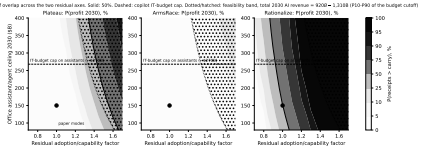
$400B

AI revenue in calendar 2028 overturns the central case

Interim marker: calendar 2027 near $330B is our path; above $350B favours the build-out.

The next data: the Q3 earnings calls of 28–30 October, Nvidia on 19 November, the Anthropic listing. The register is rebuilt after each and the paper re-issued if anything moves materially.

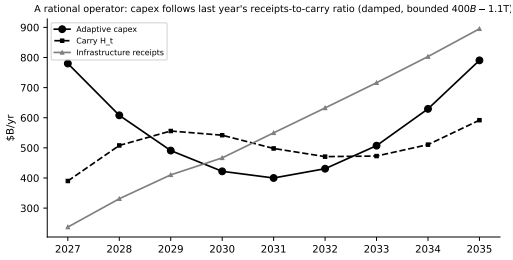
Paper Section 17 and the update protocol.



Probability of profit in 2030 over the two things we cannot pin down: how fast firms actually adopt, and how big the office-assistant market gets.

12. The way out exists, and it is not dramatic

Stop building ahead of revenue. An operator that cuts spending by about a third — or simply lets spending track the revenue it actually books — turns profitable around 2031 and recovers the hole by 2034–35. A plateau never recovers within the horizon; an arms race widens the gap every year.



Paper Appendix on the 2035 horizon. The forecast is impairments and a capex retreat in 2027–29, then a cheaper, more distributed equilibrium — not the end of AI.

The whole argument, and how to take part

In five lines.

1. The installed base costs about 28 cents per dollar per year to carry; that is $725B a year by 2030 on a flat path.
2. Revenue is $180B today and must grow 49% a year to catch it.
3. Counted sector by sector, it grows 33%, and only $480–490B of it reaches the owners.
4. The gap is $255B a year in 2030, more under the street's path, and 2027–28 is already lost.
5. Watch 2027 and 2028 revenue: $400B in 2028 proves us wrong.

Every number is yours to challenge.

The paper, the model, the data and the register of inputs (each with its evidence, its range and what would improve it) are public. Change an input, run one script, and the paper's numbers and charts regenerate; the script says whether the change is minor or forces a re-issue.

Paper and bundle: Zenodo DOI (on the paper's first page).
Model and register: the GitHub repository named there.
Comments: desmond.fitzpatrick@nvara.ai

First reviewers: George Kountouris, Francis Pedraza, Matthew Fitzpatrick, Viraj Mehta. Figures are the paper's own; no number here is newer than the paper's register.